

Frontier LLM Comparison

Capability, context, and pricing across 14 frontier models · May 2026

Color key:

1st place

2nd place

Higher % is better. Lower \$ is better.

What each benchmark actually tests

 **SWE-bench Verified**
Can the AI fix real bugs in real GitHub projects? 87% means it correctly fixed 87 out of 100 actual production bugs.

 **SWE-bench Pro**
Harder version of the bug-fixing test. Multiple programming languages, stricter rules, less chance of cheating from training data.

 **GPQA Diamond**
PhD-level questions in physics, chemistry, and biology. Tests deep scientific reasoning that can't be Googled in seconds.

 **BrowseComp**
Can the AI use a web browser to research a question and return the correct answer? Like a research assistant.

 **Terminal-Bench 2.0**
Can the AI use a Linux command line to complete tasks? Like a junior dev working in a terminal.

 **OSWorld**
Can the AI control a real desktop computer? Mouse, keyboard, multiple apps, full screens. Hardest agentic test today.

 **MCP-Atlas**
Can the AI call external tools and APIs in the right order? Critical for agentic systems and tool orchestration.

 **Context window**
How much text the AI can read at once. 1M tokens ≈ 750,000 words ≈ 8 full novels in a single prompt.

Frontier model comparison

Model	Origin	SWE-V	SWE-Pro	GPQA	Browse	Term	OSW	MCP	Context	Input \$	Output \$	License
Claude Opus 4.7	US · Anthropic	87.6	64.3	94.2	79.3	~70	78.0	77.3	200K	\$5.00	\$25.00	Closed
GPT-5.5	US · OpenAI	n/a	~58	n/a	90.1	82.7	78.7	~68	256K	\$5.00	\$15.00	Closed
GPT-5.4	US · OpenAI	~80	57.7	94.4	89.3	~75	n/a	68.1	256K	\$5.00	\$15.00	Closed
Gemini 3.1 Pro	US · Google	80.6	54.2	94.3	85.9	68	n/a	n/a	1M	\$2.00	\$12.00	Closed
DeepSeek V4 Pro	CN · DeepSeek	80.6	~55	~89	n/a	n/a	n/a	n/a	1M	\$0.27	\$1.10	MIT (open)
DeepSeek V4 Flash	CN · DeepSeek	~74	n/a	n/a	n/a	n/a	n/a	n/a	128K	\$0.07	\$0.28	MIT (open)
Kimi K2.6	CN · Moonshot	~78	n/a	n/a	n/a	n/a	n/a	n/a	256K	\$0.60	\$2.50	Mixed
GLM-5.1	CN · Zhipu	~77	58.4	n/a	n/a	n/a	n/a	n/a	200K	\$0.50	\$2.00	MIT (open)
Qwen 3.6 Max	CN · Alibaba	~79	~56	~91	n/a	n/a	n/a	n/a	1M	\$0.80	\$3.20	Closed
Qwen 3.6-35B-A3B	CN · Alibaba	73.4	n/a	86.0	n/a	n/a	n/a	n/a	256K	self-host	self-host	Apache 2.0
MiniMax M2.5	CN · MiniMax	80.2	n/a	n/a	n/a	n/a	n/a	n/a	1M	\$0.30	\$1.20	Open
MiniMax M2.7	CN · MiniMax	78.0	n/a	n/a	n/a	n/a	n/a	n/a	1M	\$0.30	\$1.20	Open
Grok 4.20	US · xAI	~76	n/a	~92	n/a	n/a	n/a	n/a	2M	\$3.00	\$15.00	Closed
Llama 4 Scout	US · Meta	~72	n/a	~84	n/a	n/a	n/a	n/a	10M	self-host	self-host	Llama 4

Notes. Prices are USD per 1 million tokens. SWE-V = SWE-bench Verified, SWE-Pro = SWE-bench Pro, GPQA = GPQA Diamond, Browse = BrowseComp, Term = Terminal-Bench 2.0, OSW = OSWorld, MCP = MCP-Atlas. **n/a** means the vendor or third-party hasn't published a verified score — not necessarily weak. Most Chinese models lack published BrowseComp / Terminal-Bench / OSWorld scores; this is a transparency gap, not a capability claim. Several non-flagship scores (~ values) are approximate from third-party evaluations. SWE-bench Verified has known training-contamination concerns affecting all frontier models — SWE-bench Pro is the more reliable successor. Sources: Anthropic, OpenAI, Google DeepMind, DeepSeek, Vellum, Artificial Analysis, Apiyi, DataCamp, Spectrum AI Lab (April–May 2026 reports).